

Introduction

This project focuses on estimating the **Take-Off Weight (TOW)** of aircraft using advanced machine learning techniques. Two primary models were developed, with one initial not used in later runs:

1. **First Model:** *Deprecated during development*
2. **Second Model (advanced.model):** Utilizes a comprehensive set of flight and weather-related features.
3. **Third Model (barebones.model):** Uses the initial challenge set, including economic indicators. Used for all values where variables for the second model could not be obtained from the trajectory files.

Both models employ **XGBoost** for regression tasks and leverage **Bayesian Optimization** for hyperparameter tuning. The models are trained on historical flight data to predict TOW with high accuracy.

Prerequisites

Libraries and Packages

Ensure that the following R packages are installed:

- `tidyverse`: Data manipulation and visualization.
- `caret`: Data preprocessing and model training utilities.
- `ParBayesianOptimization`: Performing Bayesian Optimization.
- `doParallel` and `foreach`: Parallel processing capabilities.
- `Matrix`: Handling sparse and dense matrices efficiently.
- `lubridate`: Date and time manipulation.
- `minioclient`: Interacting with MinIO object storage.
- `xgboost`: GPU-enabled version for accelerated computations.

XGBoost Installation

A GPU-enabled version of XGBoost is required for faster computation. Install it from the provided URL, which contains the pre-built GPU version compatible with your system.

System Requirements

- **GPU Support:** NVIDIA GPU with compute capability compatible with XGBoost GPU version.
- **Parallel Processing:** Multiple CPU cores to leverage parallel processing during model training and optimization.

Data Preparation

Working Directory

Set your working directory to the location where the data files are stored.

Data Loading

Second Model Data

- **Dataset:** `challenge_set_v6.csv`
- **Date Conversion:** Convert the `date` column to numeric format to facilitate modeling.

Third Model Data

- **Dataset:** Subset of `challenge_set_v6.csv` with additional economic indicators like `MSCI.Adj.Close` and `Oil.Adj.Close`.
- **Date Conversion:** Similar date conversion as in the second model.

Missing Values Handling

- **Target Variable (`tow`):** Remove any records where `tow` is missing, as these cannot be used for training the model.

Feature Engineering

Feature Selection

Second Model Features

- **Categorical Features:**
 - `aircraft_type`: Type of the aircraft.
 - `country_code_ade`: Country code of the departure airport.
 - `country_code_ades`: Country code of the arrival airport.
 - `airline`: Airline operating the flight.
 - `callsign`: Callsign of the flight.
 - `ade`: Departure airport code.
 - `ades`: Arrival airport code.
- **Numerical Features:**
 - `flight_duration`: Duration of the flight in minutes.
 - `flown_distance`: Distance flown by the aircraft in nm.
 - `ground_speed_at_lift_off`: Ground speed at the moment of lift-off.
 - `ground_speed_delta`: Difference in ground speed during take-off.
 - `taxiout_time`: Time spent taxiing out before take-off.
 - `time_to_lift_off`: Time taken from start of roll to lift-off.
 - `avg_altitude`: Average altitude during the flight.
 - `jet_stream_coeff`: Coefficient representing the effect of jet streams.
 - `date`: Numeric representation of the flight date.
 - `departure_temp`: Temperature at the departure location.
 - `arrival_temp`: Temperature at the arrival location.
 - `u_wind`: East-West component of wind speed.
 - `v_wind`: North-South component of wind speed.
 - `vertical_ascend`: Vertical ascend rate.
 - `vertical_descend`: Vertical descend rate.
 - `humidity_diff`: Difference in humidity between departure and arrival locations.

Third Model Features

- **Categorical Features:**
 - Same as in the second model.
- **Numerical Features:**
 - `flight_duration`
 - `flown_distance`
 - `taxiout_time`
 - `MSCI.Adj.Close`: Adjusted closing price of the MSCI index (economic indicator).
 - `Oil.Adj.Close`: Adjusted closing price of oil (economic indicator).
 - `date`

Data Splitting

- **Training Set:** 80% of the data.
- **Testing Set:** 20% of the data.
- **Seed:** Set a random seed (e.g., 123) for reproducibility.

Categorical Data Processing

- **One-Hot Encoding:** Convert categorical variables into binary indicator variables.
- **Sparse Matrix:** Use sparse matrices to efficiently handle high-dimensional data resulting from one-hot encoding.

Numerical Data Processing

- **Standardization:** Center and scale numerical features to have zero mean and unit variance.
- **Preprocessing Model:** Fit the preprocessing model on training data and apply it to both training and testing sets.

Feature Matrix Construction

- **Combine:** Merge the processed categorical and numerical matrices to form the final feature matrices (`X_train` and `X_test`) for training and testing.

Model Selection and Hyperparameter Optimization

Choice of XGBoost Model

XGBoost is selected for the following reasons:

- **Performance:** XGBoost is renowned for its efficiency and accuracy in regression tasks, often outperforming other algorithms.
- **Handling of Sparse Data:** Efficiently manages sparse input data, which is beneficial due to the one-hot encoding of categorical variables.
- **Regularization:** Provides built-in regularization parameters (`lambda`, `alpha`, `gamma`) to prevent overfitting.
- **Parallel Processing and GPU Support:** Supports parallel processing and GPU acceleration, significantly reducing training time for large datasets.
- **Flexibility:** Offers extensive hyperparameter options to fine-tune model performance.

Usage of Bayesian Optimization

Bayesian Optimization is used to optimize the hyperparameters of the XGBoost models.

- **Efficiency:** Bayesian Optimization is more sample-efficient than traditional methods like grid search or random search, requiring fewer iterations to find optimal hyperparameters.
- **Probabilistic Model:** Builds a probabilistic model (usually a Gaussian Process) of the objective function to make informed decisions about where to sample next.
- **Global Optimization:** Balances exploration and exploitation, efficiently searching the hyperparameter space to find the global optimum.
- **Suitability for Expensive Functions:** Ideal for optimizing functions that are expensive to evaluate, such as model training with cross-validation.

Hyperparameter Optimization Process

1. **Objective Function Definition:**
 - The objective function accepts hyperparameters as inputs.
 - It performs k-fold cross-validation using `xgb.cv` to evaluate model performance with the given hyperparameters.
 - The function returns the negative mean RMSE from cross-validation (since the optimizer maximizes the objective function).
2. **Hyperparameter Bounds:**
 - Define realistic ranges for each hyperparameter to guide the optimization process:
 - **Learning Rate (`eta`):** 0.01 to 0.5.
 - **Maximum Tree Depth (`max_depth`):** 8 to 15 (must be integers).
 - **Minimum Child Weight (`min_child_weight`):** 8 to 15.
 - **Subsample Ratio (`subsample`):** 0.5 to 0.8.
 - **Column Sample by Tree (`colsample_bytree`):** 0.5 to 0.8.
 - **Regularization Parameters:**
 - `gamma`: 0 to 10.
 - `lambda` (L2 regularization): 0 to 10.
 - `alpha` (L1 regularization): 0 to 10.
 - **Number of Boosting Rounds (`nrounds`):** 200 to 750 (must be integers).
3. **Parallel Processing Setup:**
 - Utilize available CPU cores to parallelize the optimization.
 - Create a cluster using the `doParallel` package and register it with `foreach`.
 - Export necessary variables and libraries to each worker in the cluster.
4. **Bayesian Optimization Execution:**
 - **Initialization:** Start with a predefined number of random samples (`initPoints`) to explore the hyperparameter space.
 - **Iterations:** Run the optimization for a specified number of iterations (`iters.n`).
 - **Acquisition Function:** Use an acquisition function like the Upper Confidence Bound (UCB) to decide where to sample next, balancing exploration and exploitation.
 - **Parallelization:** Execute the optimization in parallel to speed up computation.
5. **Best Hyperparameters Extraction:**
 - After optimization, extract the hyperparameters corresponding to the best (lowest) RMSE.
 - These hyperparameters are used to train the final model.

Model Training and Evaluation

Data Alignment

- **Feature Consistency:** Ensure that the training and testing datasets have identical features.
- **Handling Missing Features:** If certain features are present in the training set but absent in the testing set, add these features to the testing set with zero values to maintain alignment.

DMatrix Creation

- **Training DMatrix:** Create a DMatrix object for the training data, which is an optimized data structure that XGBoost uses for efficient computation.
- **Testing DMatrix:** Similarly, create a DMatrix object for the testing data.

Final Model Training

- **Set Parameters:** Use the best hyperparameters obtained from Bayesian Optimization.
- **Early Stopping:** Implement early stopping by monitoring the validation error and stopping training if it doesn't improve after a certain number of rounds (`early_stopping_rounds`).
- **Evaluation Metrics:** Use RMSE as the evaluation metric to measure model performance.
- **Watchlist:** Provide XGBoost with a watchlist containing the training and validation datasets to monitor performance during training.

Feature Importance Analysis

- **Importance Matrix:** Extract the feature importance scores from the trained model to understand the impact of each feature.
- **Visualization:** Plot the feature importance scores to visually interpret which features are most influential in predicting TOW.
- **Insights:** Use this information to potentially refine the feature set in future iterations.

Prediction on Submission Set

Data Preparation

- **Load Submission Data:** Read the submission dataset, ensuring that the `date` column is converted to numeric format.
- **Select Features:** Extract the same features used in the training data.
- **Preprocessing:** Apply the same one-hot encoding and standardization procedures used on the training data.
- **Feature Alignment:** Align the submission data features with the training data by adding any missing features with zero values.

Prediction

- **DMatrix Creation:** Convert the processed submission data into a DMatrix object.
- **Model Prediction:** Use the trained XGBoost models to predict TOW for each flight in the submission dataset.
- **Handling Multiple Models:** If predictions are made using both the second and third models, combine the results appropriately.

Results Preparation

- **Combine Predictions:** Merge predictions from different models if necessary, ensuring that each `flight_id` has a corresponding predicted `tow`.
- **Remove Duplicates:** Ensure there are no duplicate `flight_id` entries in the final results.
- **Prepare Submission File:** Create a DataFrame or CSV file containing the `flight_id` and the predicted `tow` values in the required format.

Saving Results

- **Output File:** Save the predictions to a CSV file (e.g., `team_bold_emu_predictions.csv`) without row names or indices.

Submission to MinIO

MinIO Setup

- **Access Credentials:** Use your provided MinIO `access_key` and `secret_key` for authentication.
- **Alias Configuration:**
 - Use the `mc alias set` command to set up an alias for the MinIO server.
 - Verify that the alias is set correctly using `mc alias ls`.

File Upload

- **Local File Path:** Specify the path to your local prediction CSV file.
- **MinIO Destination:** Define the destination bucket and object path on the MinIO server (e.g., `dc24/submissions/team_bold_emu_predictions.csv`).
- **Upload Command:** Use the `mc cp` command to upload your file to the MinIO server.
- **Verification:**
 - Check the command output or return code to confirm the upload was successful.
 - If necessary, list the contents of the destination bucket to verify the presence of your file.